




Go further, faster®



NFStest



Jorge Mora

mora@netapp.com

February 27, 2013



Outline

- Packet trace module
- Test utilities module
- Tests
- Installation and setup
- Future work



Packet trace module

- Python module
- Takes a trace file created by tcpdump as input
- Decodes one packet at a time
- Search for specific packets
- Helps create completely automated tests



Packet trace module

```
from packet.pktt import Pktt

x = Pktt("/traces/tracefile.cap")

# Iterate over all packets in the trace file
for pkt in x:
    print pkt
```



Packet trace module

```
from packet.pktt import Pktt

x = Pktt("/traces/tracefile.cap")

# Iterate over all packets in the trace file
for pkt in x:
    print pkt

x.next()
```



Packet trace module

```
from packet.pktt import Pktt

x = Pktt("/traces/tracefile.cap")

# Iterate over all packets in the trace file
for pkt in x:
    print pkt

x.next()

x.rewind([index])
```



Packet trace module

```
from packet.pktt import Pktt

x = Pktt("/traces/tracefile.cap")

# Iterate over all packets in the trace file
for pkt in x:
    print pkt

x.next()

x.rewind([index])

x.match(match_str [, maxindex])
```



Packet trace module

```
x.match(match_str [, maxindex])
```

Examples:

```
# Find the packet with both the ACK and SYN TCP flags
```

```
# set to 1
```

```
pkt = x.match("TCP.flags.ACK == 1 and TCP.flags.SYN == 1")
```

```
# Find the next NFS EXCHANGE_ID request
```

```
pkt = x.match("NFS.argop == 42")
```

```
# Find the next NFS OPEN request or reply
```

```
pkt = x.match("NFS.op == 18")
```




Packet trace module

Examples:

```
# Find all packets coming from subnet 192.168.1.0/24 using
# a regular expression
while x.match(r"IP.src == re('192\.168\.1\.\d*')"):
    print x.pkt

# Find GETATTR asking for FATTR4_FS_LAYOUT_TYPE(bit 62)
call = x.match("NFS.attr_request & (1<<62) != 0")
if call:
    # Find GETATTR reply
    xid = call.rpc.xid
    # Find reply where the number 62 is in the array
    # NFS.obj_attributes
    mstr = "RPC.xid == %d and 62 in NFS.obj_attributes" % xid
    reply = x.match(mstr)
```



Packet trace module

Examples:

```
# Find the next WRITE request
pkt = x.match("NFS.argop == 38")
if pkt:
    print pkt.nfs
    print pkt.nfs.NFSop
```

```
# Same as above, but using membership test operator instead
if ("NFS.argop == 38" in x):
    print x.pkt.nfs
    print x.pkt.nfs.NFSop
```



Packet trace module

Packet layers supported:

- ETHERNET II (RFC 894)
- IP layer (only supports v4)
- TCP layer
- RPC layer
- NFS v4.0
- NFS v4.1 including pNFS file layouts



Test utilities module

- Process command line arguments
- Functionality for PASS/FAIL
- Test grouping functionality
- Mechanism to run individual tests
- Multiple client support
- Logging mechanism
- Debug info control
- Mount/Unmount control
- Create files/directories
- Mechanism to start a packet trace
- Mechanism to capture NFS debugging messages
- Support for pNFS testing



Test utilities module

```
from nfstest.test_util import TestUtil

x = TestUtil()
x.scan_options()
x.test_group(msg)
x.test(expr, msg [, subtest][, failmsg])
x.run_tests(msg)
x.create_host(host)
x.dprint(level, msg)
x.mount()
x.umount()
x.trace_start()
x.trace_stop()
x.trace_open()
x.nfs_debug_enable(nfsdebug=flags|rpcdebug=flags)
x.nfs_debug_reset()
x.exit()
```



Test utilities module

--help

Display usage information and options available

```
$ nfstest_posix --help
```

Options:

--version	Show program's version number and exit
--help	Show this help message and exit
--file=<file>	Options file
--server=<server>	Server name or IP address
--port=<port>	NFS server port [default: 2049]
--nfsversion=3 4	NFS version [default: 4]
--minorversion=0 1	Minor version [default: 1]
--export=<export>	Exported file system to mount [default: '/']
--mtpoint=<mntpoint>	Mount point [default: '/mnt/t']
--verbose=none opts debug all bitmask	Verbose level [default: 'none']



Test utilities module

`--verbose <none|opts|info|debug|all|intbitmask>`
Verbose level for info/debug messages

```
$ nfstest_posix --server 192.168.0.11 --verbose all
```

```
$ nfstest_posix --server 192.168.0.11 --verbose 0xFF
```

```
DBG2: Mount volume: mount -t nfs4 -o minorversion=1,hard,intr ...
```

```
DBG7: Allocated aligned buffer of 8192 bytes @ 0x2349000
```

```
DBG3: Open file /mnt/t/test_dio_2557_20121011204501_f_1 for reading
```

```
DBG3: Read file 4096@258048
```

```
PASS: READ right before end of file should return correct read count (4096)
```

```
DBG3: Read file 4096@262144
```

```
PASS: READ at end of file should return read count = 0
```

```
DBG7: Freeing allocated buffers
```



Test utilities module

`--tverbose <group|normal|verbose>`

Verbose level for test messages (default: normal)

```
$ nfstest_posix --server 192.168.0.11
```

```
*** Verify POSIX API access() on NFSv4
```

```
PASS: access - file access allowed with mode F_OK
```

```
PASS: access - file access not allowed with mode F_OK for a non-existent file
```

```
PASS: access - file access allowed with mode R_OK for file with permissions 0777
```

```
PASS: access - file access allowed with mode W_OK for file with permissions 0777
```

```
PASS: access - file access allowed with mode X_OK for file with permissions 0777
```

```
...
```

```
$ nfstest_posix --server 192.168.0.11 --tverbose group
```

```
PASS: Verify POSIX API access() on NFSv4 (58 passed, 0 failed)
```




Test utilities module

`--createlog`

Create log file when specified

`--keeptraces`

Do not remove any trace files at the end of execution

`--nfsdebug <flags>`

Set NFS kernel debug flags and save log messages

`--rpcdebug <flags>`

Set RPC kernel debug flags and save log messages



Test utilities module

--bugmsgs <file>

File containing test messages to mark as bugs if they failed

--ignore

Ignore all bugs given by bugmsgs

--nfsversion 2|3|4

NFS version [default: 4]

--minorversion 0|1

NFS minorversion [default: 1]



Tests

- nfstest_pnfs
- nfstest_posix
- nfstest_delegation
- nfstest_dio
- nfstest_cache

```
$ nfstest_pnfs --server 192.168.0.11
```

*** Verify traffic for file using pNFS - READ

PASS: EXCHANGE_ID should be sent to MDS

PASS: EXCHGID4_FLAG_USE_PNFS_MDS should be set

PASS: EXCHGID4_FLAG_USE_NON_PNFS should not be set

PASS: CREATE_SESSION should be sent to MDS

PASS: SEQUENCE request should start with a sequence id of 1

PASS: RECLAIM_COMPLETE should be sent to MDS

PASS: GETATTR should be sent to MDS asking for FATTR4_LEASE_TIME

PASS: NFS server should return lease time(30) > 0

PASS: GETATTR should be sent to MDS asking for FATTR4_SUPPORTED_ATTRS

PASS: NFS server should support file type layouts (LAYOUT4_NFSV4_1_FILES)

PASS: GETATTR should be sent to MDS asking for FATTR4_FS_LAYOUT_TYPE

PASS: NFS server should return LAYOUT4_NFSV4_1_FILES in fs_layout_types

PASS: OPEN should be sent

PASS: LAYOUTGET layout type should be LAYOUT4_NFSV4_1_FILES

PASS: LAYOUTGET iomode should be LAYOUTIOMODE4_READ

PASS: LAYOUTGET should ask for full file layout

PASS: LAYOUTGET reply layout type should be LAYOUT4_NFSV4_1_FILES

...

```
$ nfstest_posix --server 192.168.0.62 --tverbose group
```

PASS: Verify POSIX API access() on NFSv3 (58 passed, 0 failed)
PASS: Verify POSIX API chdir() on NFSv3 (3 passed, 0 failed)
PASS: Verify POSIX API creat() on NFSv3 (6 passed, 0 failed)
PASS: Verify POSIX API fcntl() on NFSv3 (34 passed, 0 failed)
PASS: Verify POSIX API fdatasync() on NFSv3 (2 passed, 0 failed)
PASS: Verify POSIX API fstat() on NFSv3 (44 passed, 0 failed)
PASS: Verify POSIX API fstatvfs() on NFSv3 (20 passed, 0 failed)
PASS: Verify POSIX API fsync() on NFSv3 (2 passed, 0 failed)
PASS: Verify POSIX API link() on NFSv3 (5 passed, 0 failed)
PASS: Verify POSIX API lseek() on NFSv3 (10 passed, 0 failed)
PASS: Verify POSIX API lstat() on NFSv3 (44 passed, 0 failed)
PASS: Verify POSIX API mkdir() on NFSv3 (7 passed, 0 failed)
PASS: Verify POSIX API opendir() on NFSv3 (2 passed, 0 failed)
PASS: Verify POSIX API read() on NFSv3 (3 passed, 0 failed)
PASS: Verify POSIX API readdir() on NFSv3 (3 passed, 0 failed)
PASS: Verify POSIX API readlink() on NFSv3 (2 passed, 0 failed)
PASS: Verify POSIX API rename() on NFSv3 (10 passed, 0 failed)
PASS: Verify POSIX API rewinddir() on NFSv3 (4 passed, 0 failed)

...



Tests

```
$ nfstest_delegation --server 192.168.0.11 --client 192.168.0.18
```

*** Basic READ delegation tests

PASS: READ delegation should be granted

PASS: OPEN's should not be sent for the same file

PASS: READ stateid should be the DELEG stateid

PASS: READ's should not be sent when reading delegated file from a different process

*** Basic WRITE delegation tests

PASS: WRITE delegation should be granted

PASS: OPEN's should not be sent for the same file

PASS: WRITE stateid should be the DELEG stateid

*** Basic READ delegation tests with file lock

PASS: READ delegation should be granted

PASS: OPEN's should not be sent for the same file

PASS: READ stateid should be the DELEG stateid

PASS: READ's should not be sent when reading delegated file from a different process

PASS: LOCK should not be sent to the server

...



Tests

```
$ nfstest_dio --server 192.168.0.11
```

*** Verify eof marker is handled correctly when reading eof using aligned buffer

PASS: READ right before end of file should return correct read count (4096)

PASS: READ at end of file should return read count = 0

*** Verify data correctness when reading/writing using direct I/O

PASS: File created with buffered I/O is read correctly with direct I/O

PASS: File created with direct I/O is read correctly with buffered I/O

*** Verify fstat() gets correct file size after writing

PASS: The fstat() should get correct file size after every write

PASS: The fstat() should get correct file size after writing at offset = 10G

*** Verify READ is sent after writing when the file is open for both read & write

PASS: READ data should be correct

PASS: READ should be sent to the server

PASS: READ should be sent with correct offset (0) and count (4096)

...



Tests

```
$ nfstest_cache --server 192.168.0.11 --client 192.168.0.18 --nfsversion 3
```

*** Verify consistency of attribute caching with NFSv3 on a file acregmin = 10

PASS: File size should have not changed at t=0

PASS: File size should have not changed just before acregmin

PASS: File size should have changed just after acregmin

*** Verify consistency of attribute caching with NFSv3 on a file acregmin = 10 acregmax = 20

PASS: File size should have not changed at t=0

PASS: File size should have not changed just before acregmax

PASS: File size should have changed just after acregmax

PASS: File size should have not changed just before acregmin

PASS: File size should have changed just after acregmin

*** Verify consistency of attribute caching with NFSv3 on a directory acdirmin = 10

PASS: Hard link count should have not changed at t=0

PASS: Hard link count should have not changed just before acdirmin

PASS: Hard link count should have changed just after acdirmin

...



Installation and setup

Web page:

<http://wiki.linux-nfs.org/wiki/index.php/NFSTest>

Download and setup:

```
$ git clone git://git.linux-nfs.org/projects/mora/nfstest.git
```

```
$ cd nfstest
```

```
$ sudo python setup.py install
```

```
$ man nfstest
```



Future work

- Add option to create html output of test results
- Add Web/GUI interface
- Fix rewind() so it does not go back to the start of file
- Fix TCP retransmission issue
- Fix TCP packet ordering issue
- Add IPv6 support
- Add NFSv3 support
- Add NFSv2 support
- Create regression tests for packet trace module

Thank you

